

How I used AI in assignment

A comprehensive overview of who did what on this take-home assignment

Yagnesh L P · June 2026

PROBLEM

SOLUTION



■ Human-led ■ Bot-led, directed ■ Collaborative

1	Whiteboard the frame Worked the problem by hand first, no tool open. Sketched a rough hypothesis for Part 1 before bringing anything to a model.	HUMAN 20%
2	Seed and refine the hypothesis Took the rough hypothesis into Claude Code. Ran a five-minute check to confirm the brief's own confounds were real before ranking anything, then wrote the framing fully data-blind so it stayed pure. A MECE pass later caught one hypothesis that was quietly the cause of another and forced the whole set onto the funnel's own stage spine.	COLLABORATIVE 10%
3	Branch the exploration Ran Part 1's own five-test ladder against the data, slice by slice, instead of a naive before-and-after, branching into multiple checks at once. Alongside it, I kept a four-folder PM OS, an append-only decision log, to track findings and keep agent drift bounded.	BOT, DIRECTED 20%
4	Pressure-test the draft Drafts were built under personas, senior data analyst, product, customer, driver, then challenged in a Director-style critique pass; six objections came back and each got verified against the data rather than accepted on faith. Pushed past the brief's own example levers into a full demand-supply-match decomposition, and the first voice pass came back robotic, so I scrapped it and rewrote everything calibrated to my own published writing.	COLLABORATIVE 45%
5	Package and ship With the solution drafted end to end, Claude built the presentation layer: formatting the solution book, an AI-readable doc, structuring the repo, and a deployable page where useful.	BOT, DIRECTED 5%

The core philosophy i followed is human and AI. Each have their own strengths, and their own limitations. Key value unlock is from knowing which one the task actually needs.

WHERE I PUSHED PAST THE AI

The commission-bias catch

The model kept defaulting to Uber's playbook: commission-based pricing logic, even though NY runs zero-commission and that one fact changes what every lever actually does. It had no contextual grounding on the product or the geography, so it reached for the assumptions baked into its training instead of the brief in front of it. I caught this because I expected it going in; that kind of drift is what I watch for by default with any agent. The fix wasn't to argue it out of the assumption each time it surfaced. It was to change how I worked with it at the branching and diagnosing stages: instructed it to be collaborative rather than assertive there, surface the assumption and ask, instead of running confidently in the wrong direction.